

MASQE

AI Control Layer

Security and governance for agents that take action

HackYeah 2026



The risk begins when an agent takes action

“Review this repository.”

A legitimate request can lead the agent into instructions hidden in untrusted files.

Untrusted input


README files, retrieved documents
and tool responses

Potential damage

Stolen credentials, unsafe commands
and uncontrolled resource use

One enforcement point

Central policy and threat feeds



Rules and AI assess different risks

Deterministic controls

Who may act, on which resource?
Does the request contain a secret?
Has it exceeded its budget?

Semantic controls

Does the text try to override instructions?
Does the action drift from the task?
Does it try to poison agent memory?

Proportionate responses

Allow or redact. Route to Ghost Shell or request approval.
Throttle or block when policy requires it.

Policy changes take effect while MASQE runs

Change the configuration

Permissions, allowed models, risk thresholds and threat signatures

Control consumption

Token and cost budgets, request rates, tool calls and runtime

Invalid edits keep the last valid configuration active.

Ghost Shell: an attack without host execution

01 Read the repository

A README contains a malicious setup instruction.

02 Run the “installer”

The emulator returns a simulated result.

03 Read the fake .env

Per-session honeytokens act as marked credentials.

04 Attempt an upload

MASQE detects the tokens and records an incident.

No host command or outbound network request runs. The recorder preserves a hash chain.

A console for security and management

Security teams

Live decisions and incidents
Triggered controls and audit records
Ghost Shell command history

Management

Resource use and budget status
Latency percentiles
Exportable reports

Optional AI explanations translate the verdict into plain language.

The local explainer currently answers in Polish and runs after the security decision.

Evidence the jury can challenge

159

automated tests passed

Verified run: 4 October 2026

Positive and negative cases

Safe requests proceed. Unsafe actions trigger the configured response.

Real services, live changes

HTTP flows, policy reload, Ghost Shell, API integrations and failure handling

Try an ad-hoc prompt. Change a rule. Inspect the decision and telemetry.

MASQE today

A working local control layer

Gateway and integrations

Live policy, budgets and reporting

Ghost Shell with traceable attack evidence

Next: broader adversarial evaluation and production hardening

github.com/OskarBartoszyk/MASQE-HackYeah2026